Optimist

Primary site: yesoptimist.com. Captured 2026-10-14. Offering: agency for B2B SaaS and technology. `llms.txt` calls it “a full-service SEO & AEO agency.”

Three-line summary

Optimist publishes an operating system, not a tracker: the CORE framework with 21 factors, a fan-out study with a stated sample, and public starting prices of $2,500 for a roadmap and $3,000 per month for advisory. They benchmark ChatGPT, Claude, Perplexity, Gemini, and Google AI Overviews, and they argue SEO and AI search are one system. The site is WordPress. Yoast generates `llms.txt`.

Offering

The services page says Optimist helps B2B technology companies get cited across answer engines “building on SEO foundations” (services). `llms.txt` says consulting, advisory, and execution (llms.txt). Named relationships on the services page: Semrush, ZoomInfo, DreamHost, and HelloSign, as companies where the team has worked alongside in-house marketers. “Nearly 200 B2B tech & SaaS clients over 10 years.”

GEO methodology

CORE (“Complete Organic Revenue Engine”) scores the site and the brand across 21 factors grouped as technical foundations, content coverage, content construction, message fidelity, retrieval signals, authority, and proof. The page says the same framework has been used with nearly 200 clients.

Work products: audits, 6–12 month roadmaps, content briefs, schema and structure, entity and knowledge-graph work, digital PR, and monthly reporting. They argue against publishing more articles on topics that already have “twenty competent posts,” and for covering unanswered questions, answer-first structure, and consistent entity descriptions.

A methods claim worth isolating: “Optimist’s AI Fan-Out Study tested 199 buyer-style prompts across five AI surfaces against 1,708 Google SERPs and found that pages ranking for a model’s own follow-up queries explain its citations at roughly five times the rate of a randomized control.” That is the most specific agency experiment in this set. The data is not linked as a table in the extract. “Roughly five times” is their summary.

Pricing of the work: one-time strategy “start at $2,500.” Advisory retainers “start at $3,000 a month,” month-to-month, and include “weekly AI visibility tracking.” They also say market retainers elsewhere run “from roughly $1,500 to $25,000 a month,” which is their market color, not a census.

How they measure

Engines they say they benchmark: ChatGPT, Claude, Perplexity, Gemini, and Google AI Overviews. KPIs listed: brand recommendation rate, citation rates, share of voice versus competitors, SEO rankings, traffic from LLMs, LLM-sourced leads and pipeline, revenue influence. Weekly tracking is part of the retainer. The tracker vendor is not named. Fan-out queries are part of the audit.

Not stated: UI versus API, repeats, or a variance formula. The fan-out study is about the relationship between Google rankings and citations, not about sampler noise.

Matching real users

They say to prioritize “where your buyers already are,” and that for most B2B companies that is ChatGPT first and AI Overviews second, then to check that assumption against the client’s own benchmark in the first month. That is a locale-and-engine selection policy, not a personalization study. Five surfaces in the fan-out study match the five they report on.

Metrics

Recommendation rate, citation rate, share of voice, rankings, LLM referral traffic, leads, pipeline, revenue. Case multipliers on the page, as their claims: “26x LLM Referral Traffic,” “13x LLM Revenue,” “25x LLM Referral Traffic,” and in body copy a B2B technology client “grew LLM referral revenue 49x over 14 months” and a fintech client “grew LLM conversions 8x in eight months.” The 26x, 13x, and 25x blocks do not name the client in the extract. The 49x and 8x sentences do not name the client either. Treat them as unnamed vendor claims.

SEO overlap

The page’s thesis: “AI search optimization and SEO run as one system.” The fan-out result is the mechanism: citations follow pages that rank for the model’s own sub-queries. Cutting SEO to fund AEO “removes the foundation.” That is the strategic claim a platform’s data model should be able to test, because it is falsifiable if fan-out queries and citations are stored.

Pipeline and architecture

Facts. HAR is mostly `www.yesoptimist.com` with WordPress paths (`wp-content` 656), Google Tag Manager, Facebook, LinkedIn, Twitter/X (REA evidence `ev_6cd0cc646e20dff30d99fe7d64972739d7925815603c514dc5e31ca67709fb78`). Generator context is WordPress; `llms.txt` says “Generated by Yoast SEO v27.0.” Robots is a Yoast allow-all plus a sitemap index. Google-font CSS URLs contained `|` characters illegal in HAR 1.2; those URLs were percent-encoded before inspection.

Inferences. WordPress and Yoast are the public site. Weekly visibility tracking implies a SaaS tracker or a manual run. The fan-out study implies they stored prompts, five engine responses, the model’s follow-up queries, and the

corresponding Google results. That is a spreadsheet or a warehouse. It is not a product they sell. A platform that cannot store fan-out queries cannot reproduce their core finding.

Client claims

Nearly 200 clients, the named brand relationships, and the multipliers above. “Head of Demand Generation, ScreenMeet” and Morten Poulsen, CEO of Plytix, are quoted. The Plytix quote is about organic becoming the top channel, not about a specific LLM multiplier.

Evals and QA

The fan-out study is the eval: 199 prompts, five surfaces, 1,708 SERPs, a randomized control. No confidence interval or pre-registration is in the extract. CORE’s 21 factors are a rubric. Inter-rater reliability is not given.

Security

Not discussed. Retainers include Slack and shared planning tools. Client prompt data handling is not specified.

Pricing

Roadmap from $2,500. Advisory from $3,000 per month, month-to-month. Execution-heavy retainers are not given a higher public number. The $1,500–$25,000 range is their description of the market.

Strengths

They published a falsifiable mechanism (fan-out rankings explain citations) with a sample size. They priced the judgment layer separately from execution, which matches how a platform should be bought: the tool does not replace the roadmap. Engine list is explicit. SEO is not abandoned. `llms.txt` exists because Yoast emits it.

Weaknesses

Case multipliers are unnamed. The tracker is unnamed, so weekly “share of voice” is not reproducible by a second vendor. The fan-out study is not linked as data. Yoast `llms.txt` also emits layout-builder taxonomy URLs, which is noise for a model. No variance story.

What to copy or beat

Copy the fan-out test as a built-in report: for each citation, show the model’s sub-queries and whether the cited URL ranks for them. Copy CORE’s factor groups as an audit schema. Beat unnamed multipliers by requiring a

client, a window, and a denominator before a number can be exported. Sell the $3,000-a-month judgment as a workflow (roadmap, briefs, weekly read) inside the product, not only as an agency.

Inferences, separated

WordPress and Yoast are observed. The warehouse behind the fan-out study is inferred. Case multipliers are their claims and are not audited here.

Sources

- https://www.yesoptimist.com/ai-search-optimization-services/
- https://www.yesoptimist.com/llms.txt
- https://www.yesoptimist.com/robots.txt
- HAR and REA JSON: `rea/har/optimist.har`, `rea/json/optimist.json`

REA technical evidence

Command run on 2026-10-14, offline, against the committed capture:

```
rea inspect-web-network-capture /workspace/rea/har/optimist.har har --json
```

REA does not refetch the recorded URLs. Package observed: rea-agents 6.5.0. Decoder reported by the evidence: `REA HAR capture adapter` `har-schema@2.0.0;jsonc-parser@3.3.1;lossless-json@4.3.1` (`har-schema/lossless-json`).

- Evidence id: `ev_6cd0cc646e20dff30d99fe7d64972739d7925815603c514dc5e31ca67709fb78`
- Artifact SHA-256: `0bba6094d106c02238322148002852f146e9b3dd7b137a6be1025821d31cd411`
- Artifact bytes: 12159759
- Records: 1023
- JSON: `rea/json/optimist.json`
- HAR: `rea/har/optimist.har`

Highest host counts in the capture (request counts, not unique users): `www.yesoptimist.com` (704), `fonts.gstatic.com` (75), `www.google-analytics.com` (48), `www.googletagmanager.com` (24), `connect.facebook.net` (24), `www.google.com` (23), `staging-yesoptimist.kinsta.cloud` (14), `fonts.googleapis.com` (12).

Response `server` header samples: `cloudflare` (509), `Golfe2` (25), `Google Tag Manager` (24), `ESF` (12).

Document rows retained in the index (ordinals are the HAR entry order REA preserves):

- ordinal 0: HTTP 200 `https://www.yesoptimist.com/` (text/html; charset=UTF-8)
- ordinal 83: HTTP 404 `https://www.yesoptimist.com/pricing` (text/html; charset=UTF-8)

- ordinal 946: HTTP 200 `https://www.yesoptimist.com/robots.txt` (text/plain; charset=utf-8)
- ordinal 951: HTTP 200 `https://www.yesoptimist.com/llms.txt` (text/plain; charset=UTF-8)
- ordinal 952: HTTP 404 `https://www.yesoptimist.com/llms-full.txt` (text/html; charset=UTF-8)
- ordinal 1016: HTTP 301 `https://www.yesoptimist.com/sitemap.xml` (text/html; charset=UTF-8)
- ordinal 1017: HTTP 200 `https://www.yesoptimist.com/sitemap_index.xml` (text/xml; charset=UTF-8)

Capture limits that bind every conclusion in this report: unauthenticated public pages only; image, media, and font requests aborted; cookie and authorization headers stripped; non-document response bodies omitted; document bodies longer than 350,000 characters omitted; this environment received HTTP 403 from Otterly and Minuttia; illegal HAR 1.2 URI characters were percent-encoded on the AthenaHQ, Optimist, and Semrush captures before validation. The marketing site is not the measurement worker.